

Bank Customer Churn Prediction - EDA

Libraries dependence

```
In [1]: import pandas as pd
import numpy as np

# Matplotlib for visualization
from matplotlib import pyplot as plt
# display plots in the notebook
%matplotlib inline

# Seaborn for easier visualization
import seaborn as sns
sns.set_style('darkgrid')

# store elements as dictionary keys and their counts as dictionary values
from collections import Counter

import warnings
```

Data Source

Kaggle - Churn Modelling Classification Data Set

- This data set contains details of a bank's customers and the target variable is a binary variable reflecting the fact whether the customer left the bank (closed his account) or he continues to be a customer.
- It consists of 10,000 records with demographic and bank history information from customers from three countries, France, Germany and Spain.

```
In [2]: #Exploratory Analysis

#Read the CSV and Perform Basic Data Cleaning
```

```
In [3]: #Load the dataset
df = pd.read_csv("Bank Churn Modelling.csv")
print(f>Dataframe dimensions: {df.shape}")
df.head()
```

Dataframe dimensions: (10000, 13)

Out[3]:

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products
0	15634602	Hargrave	619	France	Female	42	2	0.00	1
1	15647311	Hill	608	Spain	Female	41	1	83807.86	1
2	15619304	Onio	502	France	Female	42	8	159660.80	3
3	15701354	Boni	699	France	Female	39	1	0.00	2
4	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerId            10000 non-null  int64
1   Surname                10000 non-null  object
2   CreditScore            10000 non-null  int64
3   Geography              10000 non-null  object
4   Gender                 10000 non-null  object
5   Age                    10000 non-null  int64
6   Tenure                  10000 non-null  int64
7   Balance                 10000 non-null  float64
8   Num Of Products        10000 non-null  int64
9   Has Credit Card        10000 non-null  int64
10  Is Active Member       10000 non-null  int64
11  Estimated Salary       10000 non-null  float64
12  Churn                   10000 non-null  int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB
```

```
In [5]: # List number of unique customer IDs
df.CustomerId.nunique()
```

Out[5]: 10000

All Cusuomer IDs are unique --> that also means no duplicates

```
In [6]: df.duplicated().sum()
```

Out[6]: 0

Unused Features

====>> To make dataframe easily readable we will drop features not needed for machine learning <<====

- CustomerId
- ~

```
In [7]: # Drop unused features
df.drop(['CustomerId', 'Surname'], axis=1, inplace=True)
print(f"Dataframe dimensions: {df.shape}")
df.head()
```

Dataframe dimensions: (10000, 11)

Out[7]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary
0	619	France	Female	42	2	0.00	1	1	1	101356
1	608	Spain	Female	41	1	83807.86	1	0	1	112588
2	502	France	Female	42	8	159660.80	3	1	0	113561
3	699	France	Female	39	1	0.00	2	0	0	93521
4	850	Spain	Female	43	2	125510.82	1	1	1	790437

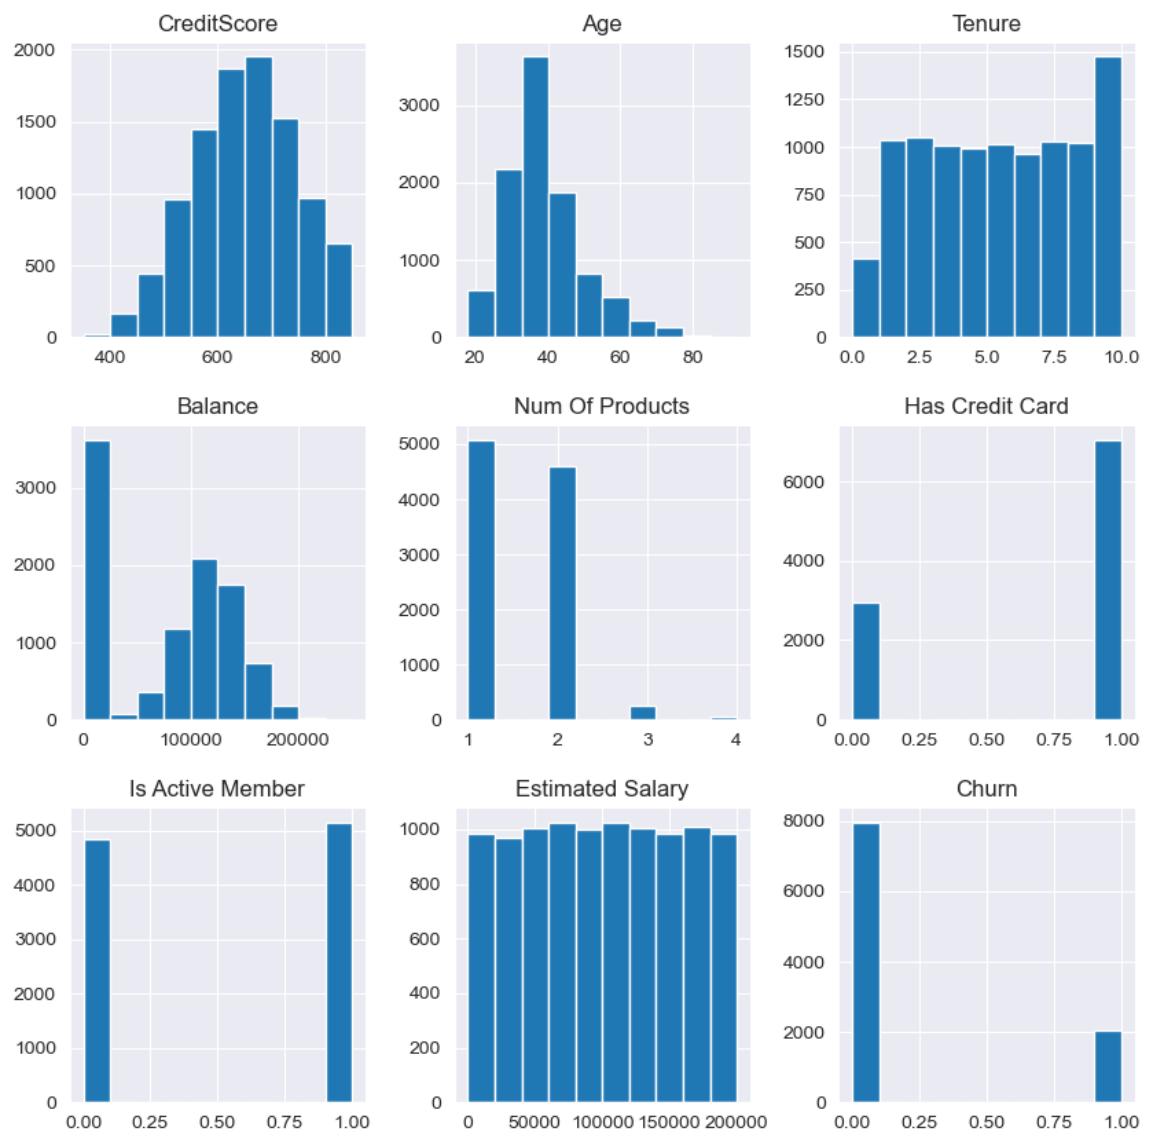
```
In [8]: df.isnull().sum()
```

```
Out[8]: CreditScore      0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
Num Of Products 0
Has Credit Card 0
Is Active Member 0
Estimated Salary 0
Churn          0
dtype: int64
```

Distributions of Numeric Features

Plotting Histogram grid

```
In [9]: # Plot histogram grid
df.hist(figsize=(10,10))
plt.show()
```



Summary statistics for the numeric features

```
In [10]: # Summarize numerical features
df.describe()
```

Out[10]:

	CreditScore	Age	Tenure	Balance	Num Of Products	Has Credit Card
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550
std	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584
min	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000
25%	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000
50%	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000
75%	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000
max	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000

From the summary statistics and the histograms we can conclude that all features look OK. We do not see any extreme values for any feature.

Distributions of Categorical Features

```
In [11]: # Summarize categorical features
df.describe(include=["object"])
```

Out[11]:

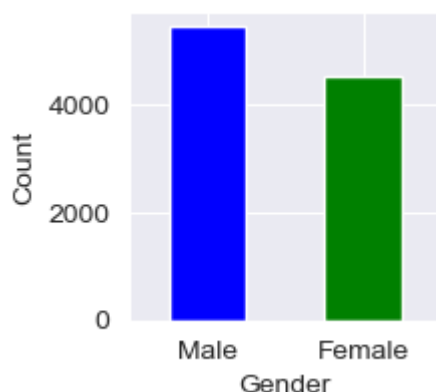
	Geography	Gender
count	10000	10000
unique	3	2
top	France	Male
freq	5014	5457

This shows us the number of unique classes for each feature. For example, there are more males (5457) than females. And France is most common of 3 geographies in our dataframe. There are no sparse classes.

Let's visualize this information

```
In [12]: # Bar plot for "Gender"
plt.figure(figsize=(2,2))
df['Gender'].value_counts().plot.bar(color=['b','g'])
plt.ylabel('Count')
plt.xlabel('Gender')
plt.xticks(rotation=0)
plt.show()

# Display count of each class
Counter(df.Gender)
```




```
Out[12]: Counter({'Female': 4543, 'Male': 5457})
```

In our data sample there are more males than females.

```
In [13]: # Bar plot for "Geography"
plt.figure(figsize=(3,3))
df['Geography'].value_counts().plot.bar(color=['b','g','r'])
plt.ylabel('Count')
plt.xlabel('Geography')
plt.xticks(rotation=0)
plt.show()

# Display count of each class
Counter(df.Geography)
```



```
Out[13]: Counter({'France': 5014, 'Spain': 2477, 'Germany': 2509})
```

Majority of customers are from France, about 50%, and from Germany and Spain around 25% each

Churn Segmentation by Gender

```
In [14]: # Segment "Churn" by gender and display the frequency and percentage within  
grouped = df.groupby('Gender')['Churn'].agg(Count='value_counts')  
grouped
```

Out[14]:

Count		
Gender	Churn	
Female	0	3404
	1	1139
Male	0	4559
	1	898

```
In [15]: # Reorganize dataframe for plotting count  
dfgc = grouped  
dfgc = dfgc.pivot_table(values = 'Count', index = 'Gender', columns = ['Churn'])  
dfgc
```

Out[15]:

Churn	0	1
Gender		
Female	3404	1139
Male	4559	898

```
In [16]: # Calculate percentage within each class
dfgp = grouped.groupby(level=[0]).apply(lambda g: round(g * 100 / g.sum(), 2))
dfgp.rename(columns={'Count' : 'Percentage'}, inplace = True)
dfgp
```

C:\Users\baps\AppData\Local\Temp\ipykernel_21152\4234812768.py:2: FutureWarning: Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included in the index, regardless of whether the applied function returns a like-indexed object. To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
dfgp = grouped.groupby(level=[0]).apply(lambda g: round(g * 100 / g.sum(), 2))
```

Out[16]:

Percentage		
Gender	Churn	
Female	0	74.93
	1	25.07
Male	0	83.54
	1	16.46

```
In [17]: # Reorganize dataframe for plotting percentage
dfgp = dfgp.pivot_table(values='Percentage', index='Gender', columns=['Churn'])
dfgp
```

Out[17]:

Churn		
	0	1
Gender		
Female	74.93	25.07
Male	83.54	16.46


```
In [18]: # Churn distribution by gender, count + percentage
```


```
labels= ['Stays', 'Churn']

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))

dfgc.plot(kind='bar',
           color=['g', 'r'],
           rot=0,
           ax=ax1)
ax1.legend(labels)
ax1.set_title('Churn Risk per Gender (Count)', fontsize=14, pad=10)
ax1.set_ylabel('Count',size=12)
ax1.set_xlabel('Gender', size=12)

dfgp.plot(kind='bar',
           color=['g', 'r'],
           rot=0,
           ax=ax2)
ax2.legend(labels)
ax2.set_title('Churn Risk per Gender (Percentage)', fontsize=14, pad=10)
ax2.set_ylabel('Percentage',size=12)
ax2.set_xlabel('Gender', size=12)

plt.show()
```



In percentage females are more likely to leave the bank; 25% comparing to males, 16%.

Churn Segmentation by Geography

```
In [19]: # Segment "Exited" by geography and display the frequency and percentage with  
grouped = df.groupby('Geography')['Churn'].agg(Count='value_counts')  
grouped
```

Out[19]:

		Count	
Geography	Churn		
France	0	4204	
	1	810	
Germany	0	1695	
	1	814	
Spain	0	2064	
	1	413	

```
In [20]: # Reorganize dataframe for plotting count  
dfgeoc = grouped  
dfgeoc = dfgeoc.pivot_table(values='Count', index='Geography', columns=['Churn'])  
dfgeoc
```

Out[20]:

		Churn	0	1
		Geography		
		France	4204	810
		Germany	1695	814
		Spain	2064	413

```
In [21]: # Calculate percentage within each class
dfgeop = grouped.groupby(level=[0]).apply(lambda g: round(g * 100 / g.sum()),
dfgeop.rename(columns={'Count': 'Percentage'}, inplace=True)
dfgeop
```

C:\Users\baps\AppData\Local\Temp\ipykernel_21152\11712713.py:2: FutureWarning: Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included in the index, regardless of whether the applied function returns a like-indexed object.
To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
dfgeop = grouped.groupby(level=[0]).apply(lambda g: round(g * 100 / g.sum(), 2))
```

Out[21]:

		Percentage
Geography	Churn	
France	0	83.85
	1	16.15
Germany	0	67.56
	1	32.44
Spain	0	83.33
	1	16.67

```
In [22]: # Reorganize dataframe for plotting percentage
dfgeop = dfgeop.pivot_table(values='Percentage', index='Geography', columns=
dfgeop
```

Out[22]:

		Churn	0	1
		Geography		
France			83.85	16.15
Germany			67.56	32.44
Spain			83.33	16.67

```

In [23]: # Churn distribution by geography, count + percentage

labels= ['Stays', 'Churn']


fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))

dfgeoc.plot(kind='bar',
             color=['g', 'r'],
             rot=0,
             ax=ax1)
ax1.legend(labels)
ax1.set_title('Churn Risk per Geography (Count)', fontsize=14, pad=10)
ax1.set_ylabel('Count',size=12)
ax1.set_xlabel('Geography', size=12)

dfgeop.plot(kind='bar',
             color=['g', 'r'],
             rot=0,
             ax=ax2)
ax2.legend(labels)
ax2.set_title('Churn Risk per Geography (Percentage)', fontsize=14, pad=10)
ax2.set_ylabel('Percentage',size=12)
ax2.set_xlabel('Geography', size=12)

plt.show()

```



The smallest number of customers are from Germany but it looks that they are most likely to leave the bank. Almost one third of German customers in our sample left the bank

Correlations

```
In [24]: # Calculate correlations between numeric features
correlations = df.corr()

# sort features in order of their correlation with "Exited"
sort_corr_cols = correlations.Churn.sort_values(ascending=False).keys()
sort_corr = correlations.loc[sort_corr_cols, sort_corr_cols]
sort_corr
```

C:\Users\baps\AppData\Local\Temp\ipykernel_21152\749194107.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlations = df.corr()
```

Out[24]:

	Churn	Age	Balance	Estimated Salary	Has Credit Card	Tenure	CreditScore	Num Of Products	Is Active Member
Churn	1.000000	0.285323	0.118533	0.012097	-0.007138	-0.014001	-0.027094	-0.047820	-0.156128
Age	0.285323	1.000000	0.028308	-0.007201	-0.011721	-0.009997	-0.003965	-0.030680	0.085472
Balance	0.118533	0.028308	1.000000	0.012797	-0.014858	-0.012254	0.006268	-0.304180	-0.010084
Estimated Salary	0.012097	-0.007201	0.012797	1.000000	-0.009933	0.007784	-0.001384	0.014204	-0.011421
Has Credit Card	-0.007138	-0.011721	-0.014858	-0.009933	1.000000	0.022583	-0.005458	0.003183	-0.011866
Tenure	-0.014001	-0.009997	-0.012254	0.007784	0.022583	1.000000	0.000842	0.013444	-0.028362
CreditScore	-0.027094	-0.003965	0.006268	-0.001384	-0.005458	0.000842	1.000000	0.012238	0.025651
Num Of Products	-0.047820	-0.030680	-0.304180	0.014204	0.003183	0.013444	0.012238	1.000000	0.000000
Is Active Member	-0.156128	0.085472	-0.010084	-0.011421	-0.011866	-0.028362	0.025651	0.000000	1.000000


Let's use Seaborn's .heatmap() function to visualize the correlation grid.

```
In [25]: # Generate a mask for the upper triangle
corr_mask = np.zeros_like(correlations)
corr_mask[np.triu_indices_from(corr_mask)] = 1

# Make the figsize 6X6
plt.figure(figsize = (6,6))

# Plot heatmap of annotated correlations; change background to white
##with sns.axes_style('white'):
sns.heatmap(sort_corr*100,
            cmap='RdBu',
            annot=True,
            fmt='.0f',
            mask=corr_mask,
            cbar=False)

plt.title('Correaltions by Churn', fontsize=12)
plt.yticks(rotation=0)
plt.show()
```



Very weak correlations in general. Only weak positive correlation with age, very weak positive correlation with balance, and very weak negative correlations with number of products and membership

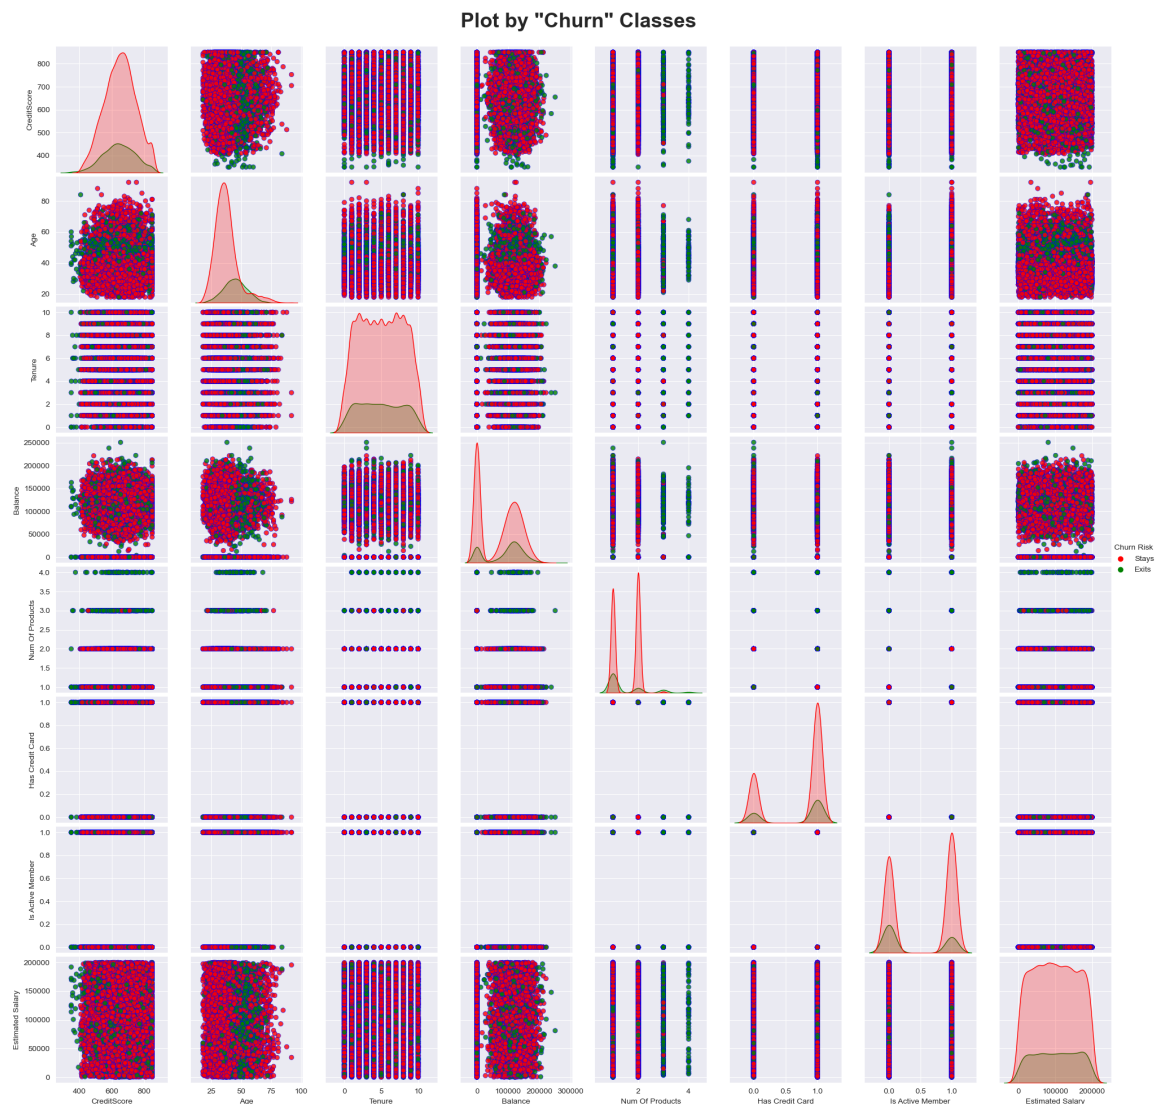
Pairplot

```
In [26]: # Plot Seaborn's pairplot
g = sns.pairplot(df, hue='Churn',
                 palette={1 : 'green',
                          0 : 'red'},
                 plot_kws={'alpha' : 0.8, 'edgecolor' : 'b', 'linewidth' : 0.5})

fig = g.fig
fig.subplots_adjust(top=0.95, wspace=0.2)
fig.suptitle('Plot by "Churn" Classes',
             fontsize=26,
             fontweight='bold')

# Update the Legend
new_title = 'Churn Risk'
g._legend.set_title(new_title)
# replace labels
new_labels = ['Stays', 'Exits']
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)

plt.show()
```



The density plots on the diagonal make it easier to compare these distributions. We can notice that only few features have slightly different distributions. For example, from the density plot for Age, it could be seen that older people have slightly higher tendency to leave the bank.

Let's reduce the clutter by plotting only four features:

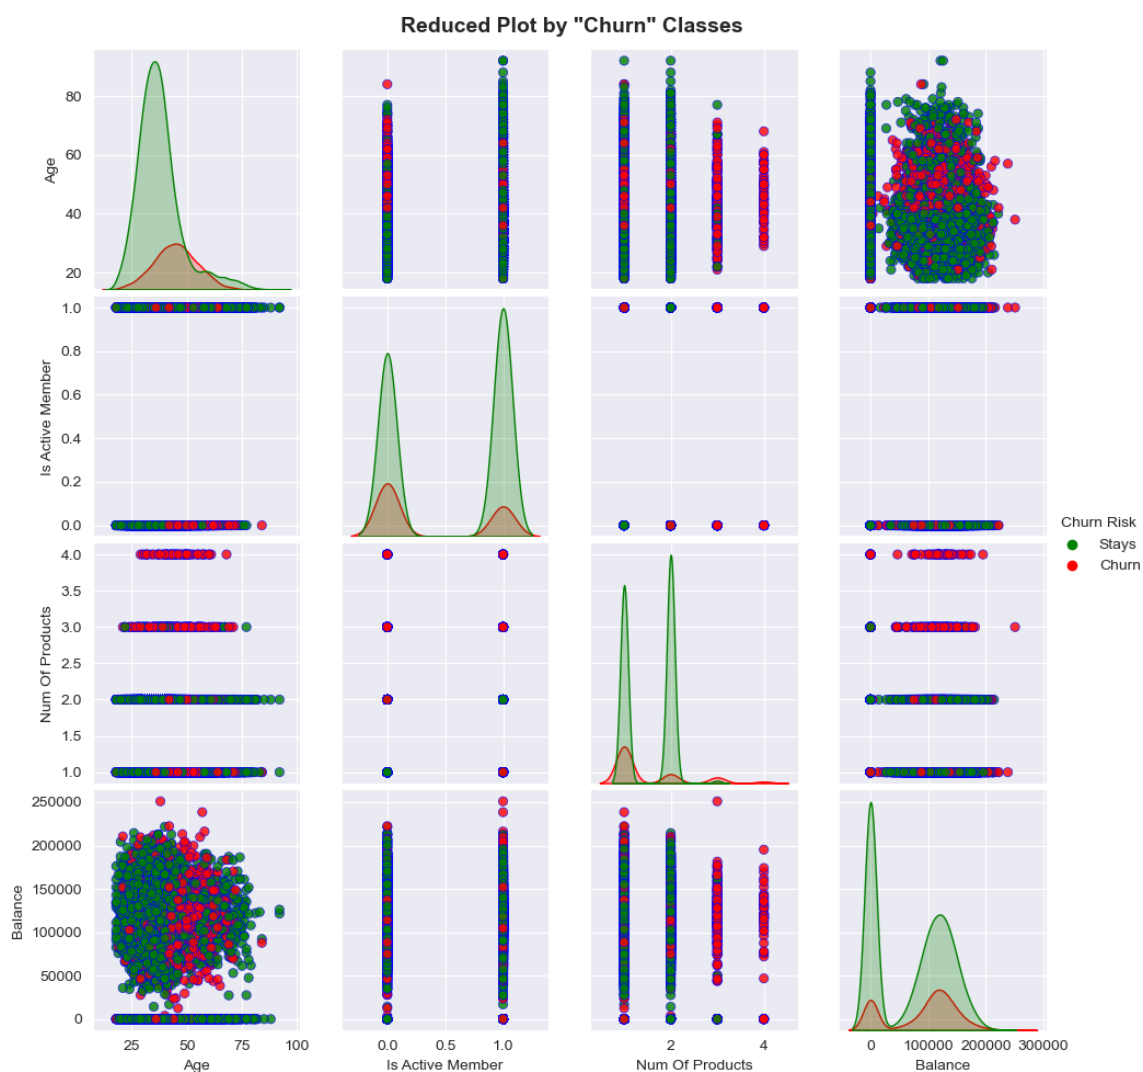
- Age,
- Is Active Member,
- Num Of Products
- Balance


```
In [27]: # Plot Seaborn's pairplot
g = sns.pairplot(df, hue='Churn',
                 vars=['Age', 'Is Active Member', 'Num Of Products', 'Balance'],
                 palette={0 : 'green',
                          1 : 'red'},
                 plot_kws={'alpha' : 0.8, 'edgecolor' : 'b', 'linewidth' : 0.5})

fig = g.fig
fig.subplots_adjust(top=0.95, wspace=0.2)
fig.suptitle('Reduced Plot by "Churn" Classes',
             fontsize=14,
             fontweight='bold')

# Update the Legend
new_title = 'Churn Risk'
g._legend.set_title(new_title)
# replace labels
new_labels = ['Stays', 'Churn']
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)

plt.show()
```



Violin Plots

```
In [28]: # Segment age by Churn and plot distributions
# "categorical" variable Churn is a numeric
# for plotting purposes only we will change it to real categorical variable

# Define palette
my_pal = {'Stays': 'green', 'Churn': 'red'}
# Convert to categorical
hr = {0: 'Stays', 1: 'Churn'}
churn = df['Churn'].map(hr)

fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(10, 8))
fig.suptitle('Churn Risk vs. Different Attributes', fontsize=14)
fig.subplots_adjust(top=0.92, wspace=0.3, hspace=0.3)

sns.violinplot(x=churn,
               y=df['Age'],
               order=['Stays', 'Churn'],
               palette=my_pal,
               ax=ax1)

ax1.set_title('Churn Risk vs. Age', fontsize=12, pad=10)
ax1.set_ylabel('Age', size=10)
ax1.set_xlabel('Churn Risk ("Churn")', size=10)

sns.violinplot(x=churn,
               y=df['Balance'],
               order=['Stays', 'Churn'],
               palette=my_pal,
               ax=ax2)

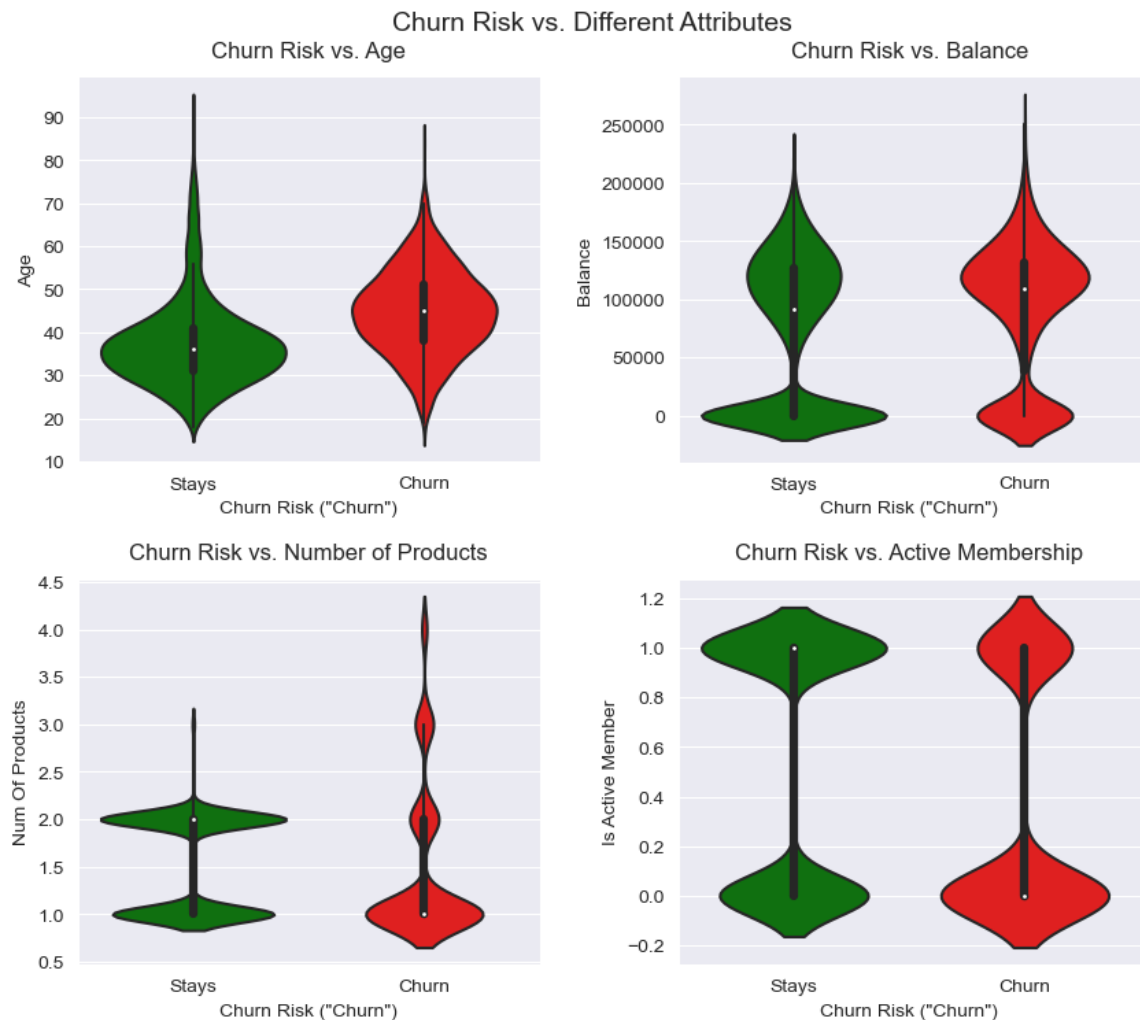
ax2.set_title('Churn Risk vs. Balance', fontsize=12, pad=10)
ax2.set_ylabel('Balance', size=10)
ax2.set_xlabel('Churn Risk ("Churn")', size=10)

sns.violinplot(x=churn,
               y=df['Num Of Products'],
               order=['Stays', 'Churn'],
               palette=my_pal,
               ax=ax3)

ax3.set_title('Churn Risk vs. Number of Products', fontsize=12, pad=10)
ax3.set_ylabel('Num Of Products', size=10)
ax3.set_xlabel('Churn Risk ("Churn")', size=10)

sns.violinplot(x=churn,
               y=df['Is Active Member'],
               order=['Stays', 'Churn'],
               palette=my_pal,
               ax=ax4)

ax4.set_title('Churn Risk vs. Active Membership', fontsize=12, pad=10)
ax4.set_ylabel('Is Active Member', size=10)
ax4.set_xlabel('Churn Risk ("Churn")', size=10)
plt.show()
```



Violin plots are confirming the earlier statement that older customers and customer with more products are more likely to leave the bank.

```
In [29]: # Define our target variable
y = df.Churn
```

```
In [30]: y.shape
```

```
Out[30]: (10000,)
```

Let's define a small helper function which displays count and percentage per class of the target feature.

```
In [31]: # Function to display count and percentage per class of target feature
def class_count(a):
    counter=Counter(a)
    kv=[list(counter.keys()),list(counter.values())]
    dff = pd.DataFrame(np.array(kv).T, columns=['Churn','Count'])
    dff['Count'] = dff['Count'].astype('int64')
    dff['%'] = round(dff['Count'] / a.shape[0] * 100, 2)
    return dff.sort_values('Count',ascending=False)
```

In [32]: *# Let's use the function*

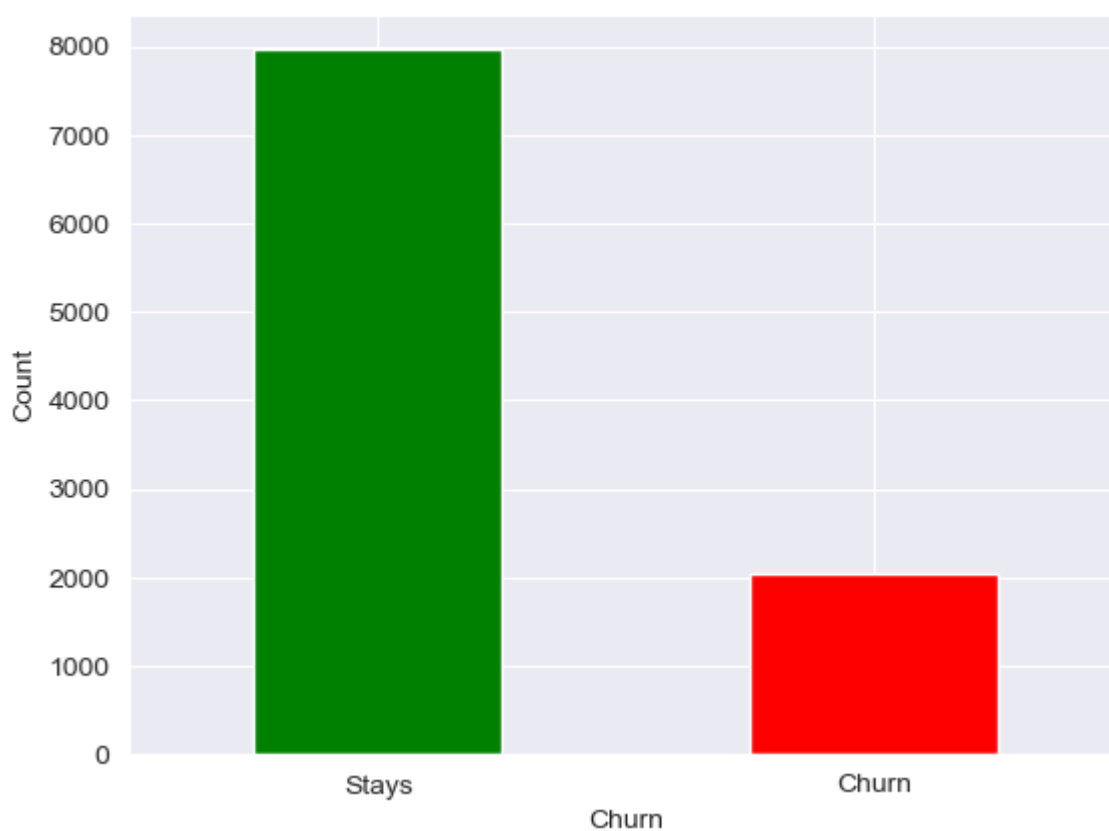
```
dfcc = class_count(y)
dfcc
```

Out[32]:

	Churn	Count	%
1	0	7963	79.63
0	1	2037	20.37

In [33]: *# Plot distribution of target variable, Exited column*

```
labels=['Stays', 'Churn']
dfcc.plot.bar(x='Churn', y='Count', color=['g', 'r'], legend=False)
plt.xticks(dfcc['Churn'], labels, rotation=0)
plt.ylabel('Count')
plt.show()
```



We can see that our dataset is imbalanced. The majority class, "Stays" (0), has around 80% data points and the minority class, "Churn" (1), has around 20% datapoints.

To address this, in our machine learning algorithms we will use SMOTE (Synthetic Minority Over-sampling Technique).

Finalizing the Dataframe

```
In [34]: df.head()
```

Out[34]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary
0	619	France	Female	42	2	0.00	1	1	1	1013
1	608	Spain	Female	41	1	83807.86	1	0	1	1125
2	502	France	Female	42	8	159660.80	3	1	0	1139
3	699	France	Female	39	1	0.00	2	0	0	938
4	850	Spain	Female	43	2	125510.82	1	1	1	790

```
In [35]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   CreditScore            10000 non-null  int64  
1   Geography              10000 non-null  object  
2   Gender                 10000 non-null  object  
3   Age                   10000 non-null  int64  
4   Tenure                 10000 non-null  int64  
5   Balance                10000 non-null  float64 
6   Num Of Products        10000 non-null  int64  
7   Has Credit Card        10000 non-null  int64  
8   Is Active Member       10000 non-null  int64  
9   Estimated Salary       10000 non-null  float64 
10  Churn                  10000 non-null  int64  
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

Our dataframe looks good and it is ready to be saved.

Save the dataframe as the analytical base table

```
In [36]: # Save analytical base table
        ### df.to_csv('ELCOT\Desktop\csv file_base_table.csv', index=None)
```